

IDR Reconciliation Dashboard

Identity resolution monitoring across EDL and CDP (Kafka) — what each tab shows, the metrics behind it, and how to read it.

Audience Analysts through senior leadership

Grain Daily, by event date

Refresh Weekly

Platform Microsoft Power BI

Why this dashboard exists

Identity resolution (IDR) stitches a login into a known customer profile. It succeeds only when three events line up for the same visit — the **Kafka login event**, the **web or app event**, and the **account event**. When any one of them is missing, the profile does not resolve, and that customer becomes invisible to everything downstream that depends on identity.

At the same time, the same login traffic is recorded on two platforms: **EDL** and **CDP (Kafka)**. If those two disagree on volume, the resolution rate itself is being measured against an incomplete picture.

This dashboard answers three questions, in order:

1. **Do the two platforms agree on what happened?** — volume reconciliation
2. **How much of it resolved?** — the headline rate
3. **For what did not resolve, which missing event caused it?** — the actionable part

Data is held at **daily event-date grain**, and every page responds to one shared date-range filter.

The headline number

Overall IDR Resolution Rate (%)

`sum(IDR Resolved) / sum(IDR Resolved + IDR Not Resolved)`

Calculated across every date selected in the filter.

This is deliberately a **volume-weighted period rate**, not an average of daily rates. A low-traffic day cannot pull the number around the way a simple daily average would, so the figure on screen always means: *of all profiles attempted in this window, this share resolved*.

Two things follow. The number **changes when you change the date range** — that is by design, not a defect. And it can always be checked by hand against **Profiles Attempted**, which is shown alongside it.

The seven pages

1 · Overview

THE FIVE-SECOND READ — SENIOR LEADERSHIP, STANDING REVIEWS

The health of the whole pipeline on one screen: overall resolution rate, non-resolution rate, profiles attempted, session-level and profile-level match rates between EDL and CDP, the **IDR Platform Health Score (0–100)**, the count of active alerts, and the single largest reason resolution is currently failing.

Every tile is colour-scored against **defined thresholds**, not against a gut feel — so “amber” means the same thing every week.

If nothing is amber or red here, there is nothing to do. If something is, the remaining pages exist to explain it.

2 · EDL vs CDP Volume

DO THE PLATFORMS AGREE? — DATA ENGINEERING, PLATFORM OWNERS

Volume reconciliation at both grains:

- **Session level (sessionid)** — Total EDL Logins vs Total Kafka Login Events
- **Profile level (syfid)** — Unique EDL Logins vs Unique Kafka Logins at Profile

Plus the signed differences and both match rates as trends, and a **device-level breakdown comparing EDL against CDP separately for web and for app** — so a gap can be attributed to a channel rather than left as an aggregate.

Reading it. 100% means the platforms agree. A persistent gap at *session* level points at events being lost in transit. A gap that appears only at *profile* level points at a de-duplication difference instead. Those are different teams and different fixes, which is why both are shown.

3 · IDR Resolution

THE RATE, IN FULL — ANALYSTS, PRODUCT OWNERS

The headline rate broken out over time and by channel: Web IDR Resolved and App IDR Resolved, their shares of the total, and the EDL-side counterparts (`web_idr_resolved_edl` and `apps_idr_resolved_edl` — resolution measured at syfid level *in EDL*).

This is where you establish whether a movement in the rate is web-driven or app-driven, and whether EDL and CDP tell the same story about it.

4 · Why Resolution Failed

THE PAGE THAT DRIVES ACTION — WHOEVER IS BEING ASKED TO FIX IT

Non-resolution is decomposed into **eight partial-match populations**: Web Only, App Only, Kafka Only, Kafka + Web, Kafka + Apps, Web + Account, Kafka + Account, and Apps + Account.

These are not estimates. They **sum exactly to IDR Not Resolved**, so every unresolved profile has a stated reason and nothing is left unexplained.

Each population is ranked by share, trended weekly, and — most usefully — **grouped by the event that was missing**. Eight raw counts become one instruction: *this is the event to go and fix*.

That grouping surfaces as **Top Missing Event** and **Top Non-Resolution Reason**, which also appear on the Overview, so the escalation path is visible from the front page.

5 · Daily Detail

THE AUDIT TRAIL — ANYONE CHECKING A SPECIFIC DATE

Every source field at daily grain, using the workbook's own column names. This is the page to open when someone asks "what happened on the 14th," or wants to reconcile a figure against the source file. Nothing here is transformed or rounded away.

6 · Alerts & Data Quality

WHAT THE REPORT IS NOT SURE ABOUT — REPORT OWNERS, DATA STEWARDS

The alert engine evaluates named rules against thresholds and rolling baselines, reporting each with a severity, its current value, its baseline, the change, and a written message. Alerts fire on breach — they do not sit there as noise the rest of the time.

The data-quality layer checks incoming data for nulls, out-of-range values, duplicate dates, calendar gaps, and internal contradictions — including whether the resolution rate published in the source file agrees with the rate this report calculates.

Governing principle: this layer **flags, it never silently corrects**. If the source data disagrees with itself, you will see the disagreement rather than a quietly patched number. Trusting the top-line rate means being able to see when it should not be trusted.

7 · Definitions

NO AMBIGUITY ABOUT WHAT A WORD MEANS — EVERYONE

A full glossary. Source fields carry **the business's own definitions, read verbatim from the workbook's dictionary tab** — nothing is paraphrased or "improved." Calculated metrics are listed separately, each with its formula, business meaning, threshold, and the pages that use it.

If two people disagree about what a number means, this page settles it.

Design flow

The pages are ordered as a funnel, and that order is the intended reading path:

[Is anything wrong?](#) Overview → [Is the input trustworthy?](#) EDL vs CDP Volume → [How big is it?](#) IDR Resolution → [What is causing it?](#) Why Resolution Failed → [Show me the rows.](#) Daily Detail → [What does the report itself flag?](#) Alerts & Data Quality → [What does this term mean?](#) Definitions

A leader can stop after page 1. An analyst can go all the way to page 5. Nobody has to start at page 5 to find out whether page 1 was fine.

Conventions worth knowing

- **One date filter, synced across every page.** Set the range once — it follows you as you navigate. Changing it changes every number, including the headline rate.
- **“As of” is the latest date in the data,** never today’s date. If the weekly load has not run, the dashboard says so rather than quietly presenting a stale figure as current.
- **Thresholds are configuration, not code.** Amber and red boundaries live in a table and can be retuned without rebuilding the report.
- **Field names match the source workbook.** A name on screen can be found in the Excel file without translation.
- **Refresh cadence is weekly,** by appending new rows to the bottom of the `idr` tab. The calendar, the slicer range, the rolling baselines and the trends all extend themselves — nothing needs to be re-pointed.

Where to start, by role

IF YOU ARE...	START AT	THE QUESTION YOU ARE ANSWERING
Senior leadership	Overview	Is the health score in range, and are there critical alerts?
Product / analytics	IDR Resolution, then Why Resolution Failed	Is the rate moving, and which population is driving it?
Data engineering	EDL vs CDP Volume	Do the platforms agree, and at which grain do they stop agreeing?
Report owner	Alerts & Data Quality	Is the underlying data sound enough to act on?

Questions about a metric definition should go to the Definitions page first — it carries the authoritative wording for every field.